



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Advanced Natural Language Processing

Lecture 8: Pretrained Language Model



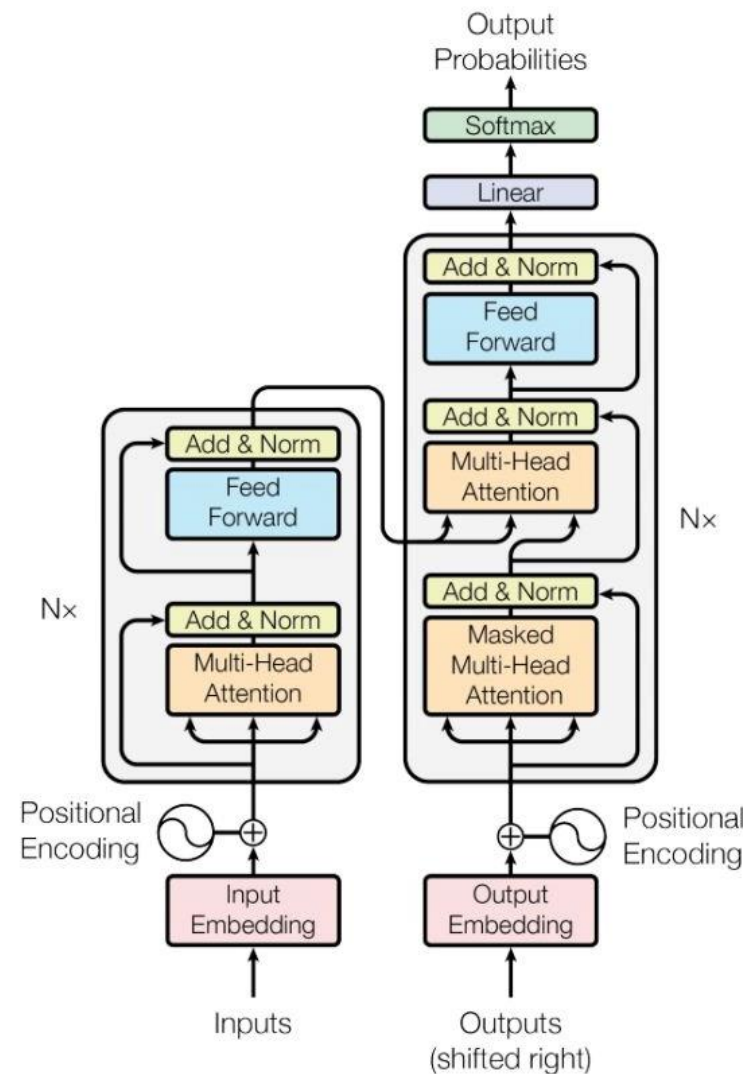
陈冠华 CHEN Guanhua

Department of Statistics and Data Science

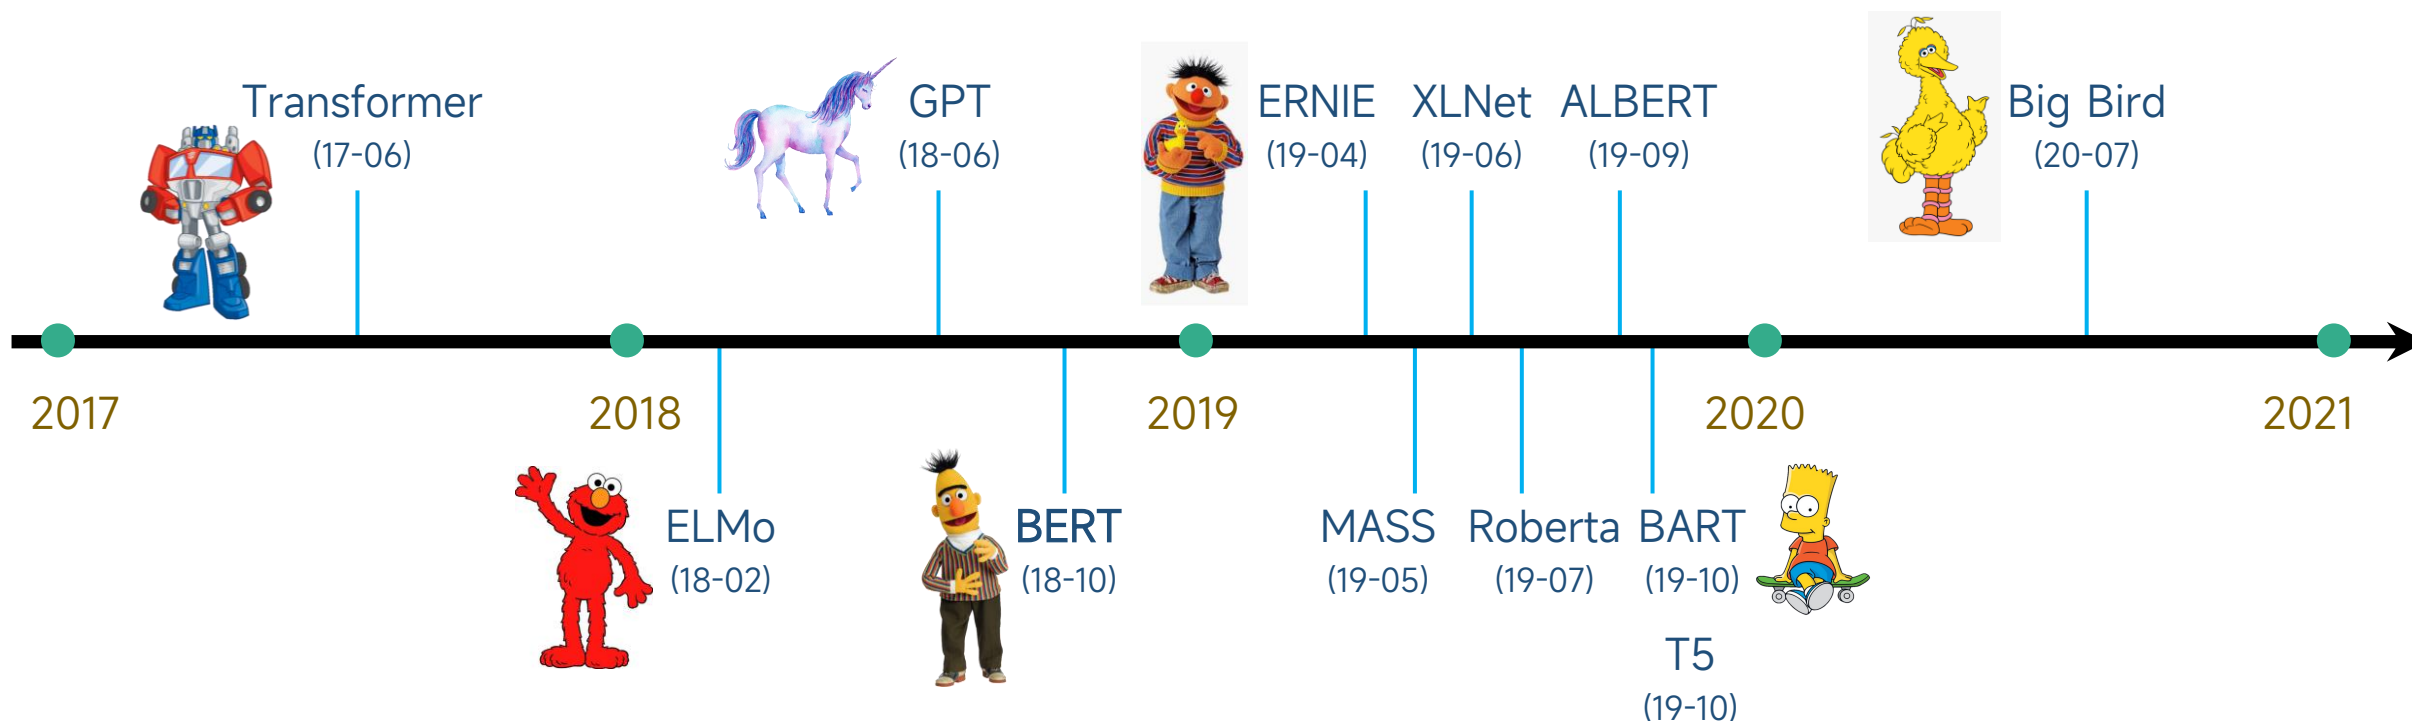
Pretrained Language Model



- Pretraining
 - The model is trained with unlabeled data in a self-supervised manner
 - E.g., masked language modeling, next token prediction
- Finetuning
 - Initialized with the pretrained model and trained with the supervised dataset of the target/downstream task



Timeline of Pretrained Language Models



Different Types of Pretrained Models



Encoder-only



Decoder-only

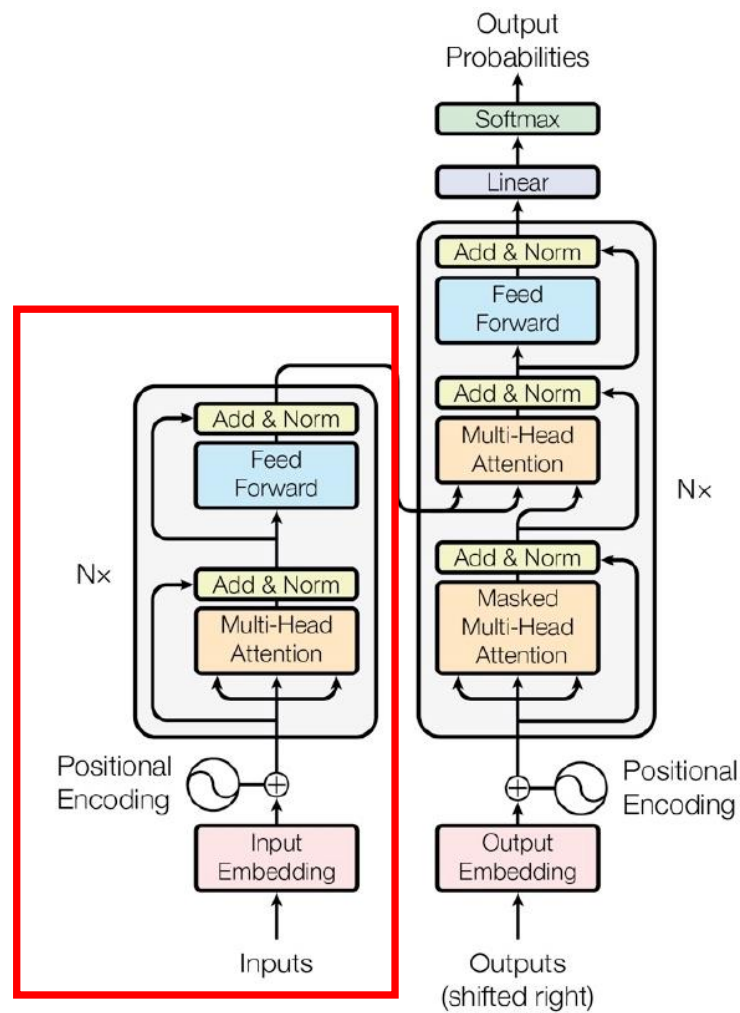


Encoder-Decoder

Pretrained Encoder



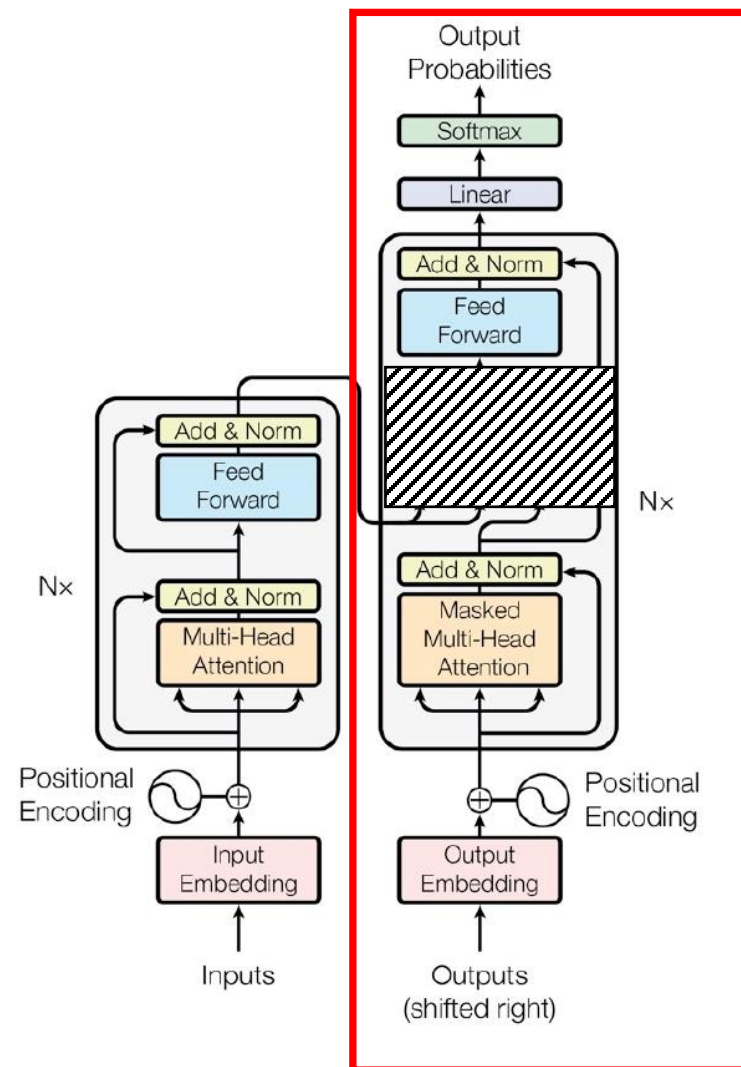
Encoder-only



Pretrained Decoder



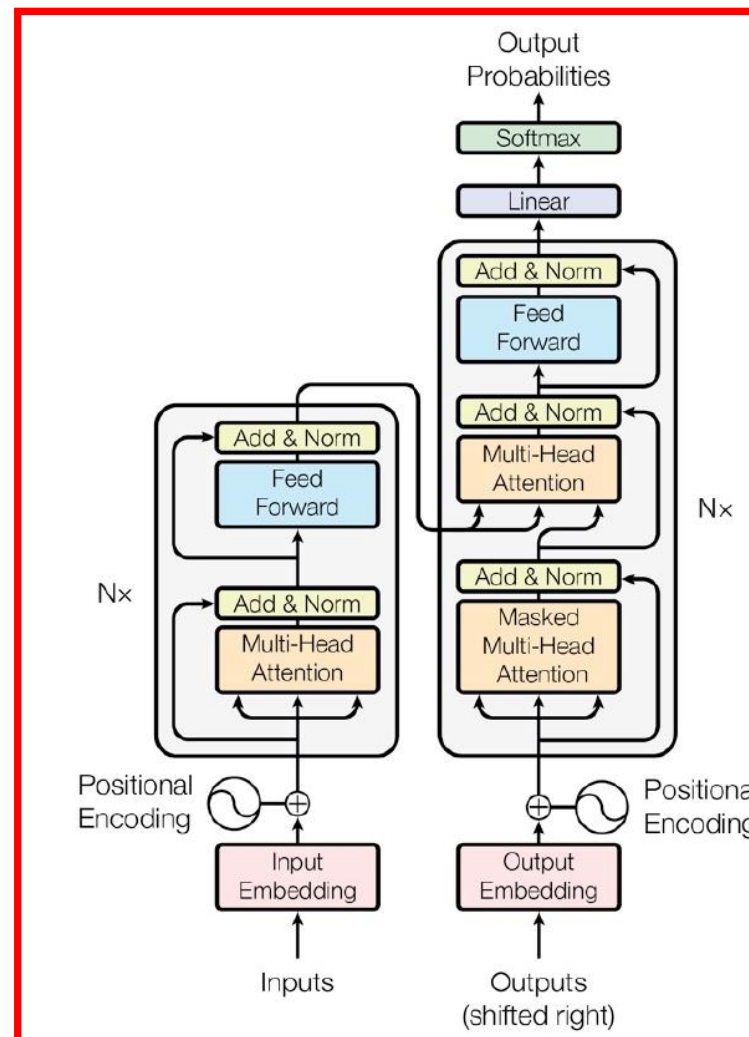
Decoder-only



Pretrained Encoder-Decoder




Encoder-decoder




Self-Supervised Learning

Supervised learning



Self-supervised learning



Prior Work: ELMo




ELMo (Peters et al., 2018; NAACL 2018 best paper)


- Train two separate **unidirectional** LMs (left-to-right and right-to-left) based on **LSTMs**
- **Feature-based** approach: pre-trained representations used as input to task-specific models
- Trained on **single sentences** from 1B word benchmark (Chelba et al., 2014)



Train Separate Left-to-Right and Right-to-Left LMs



Apply as “Pre-trained Embeddings”




Prior Work: OpenAI GPT




OpenAI GPT (Radford et al., 2018; released in 2018/6)

- Train one unidirectional LM (left-to-right) based on a deep **Transformer decoder**
- **Fine-tuning** approach: all pre-trained parameters are re-used & updated on downstream tasks
- Trained on 512-token segments on BooksCorpus — much **longer** context!


Train Deep (12-layer) Transformer LM



Fine-tune on Classification Task

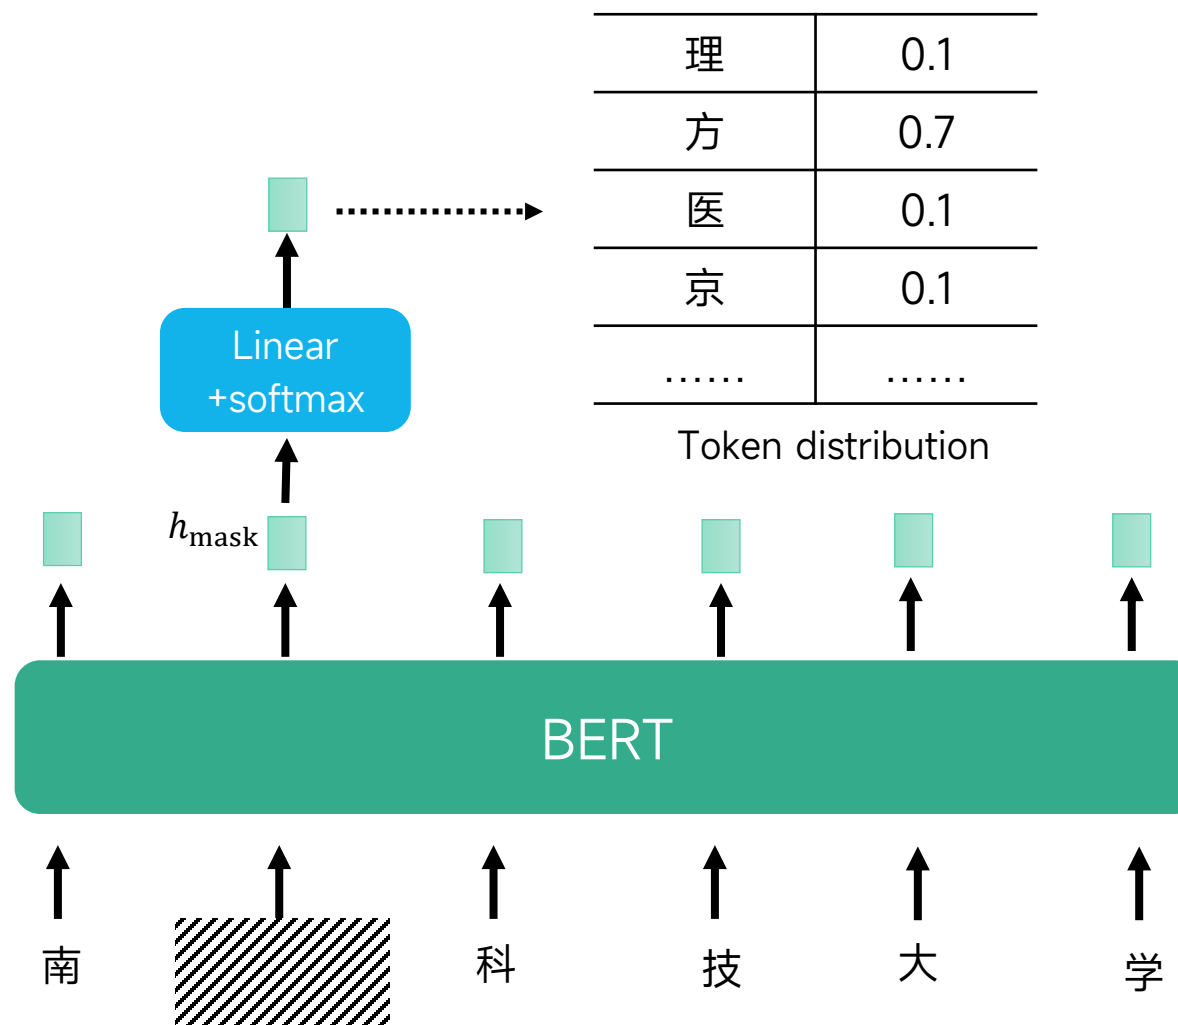


Masked Token Prediction




Transformer Encoder

Input x' (Randomly mask)

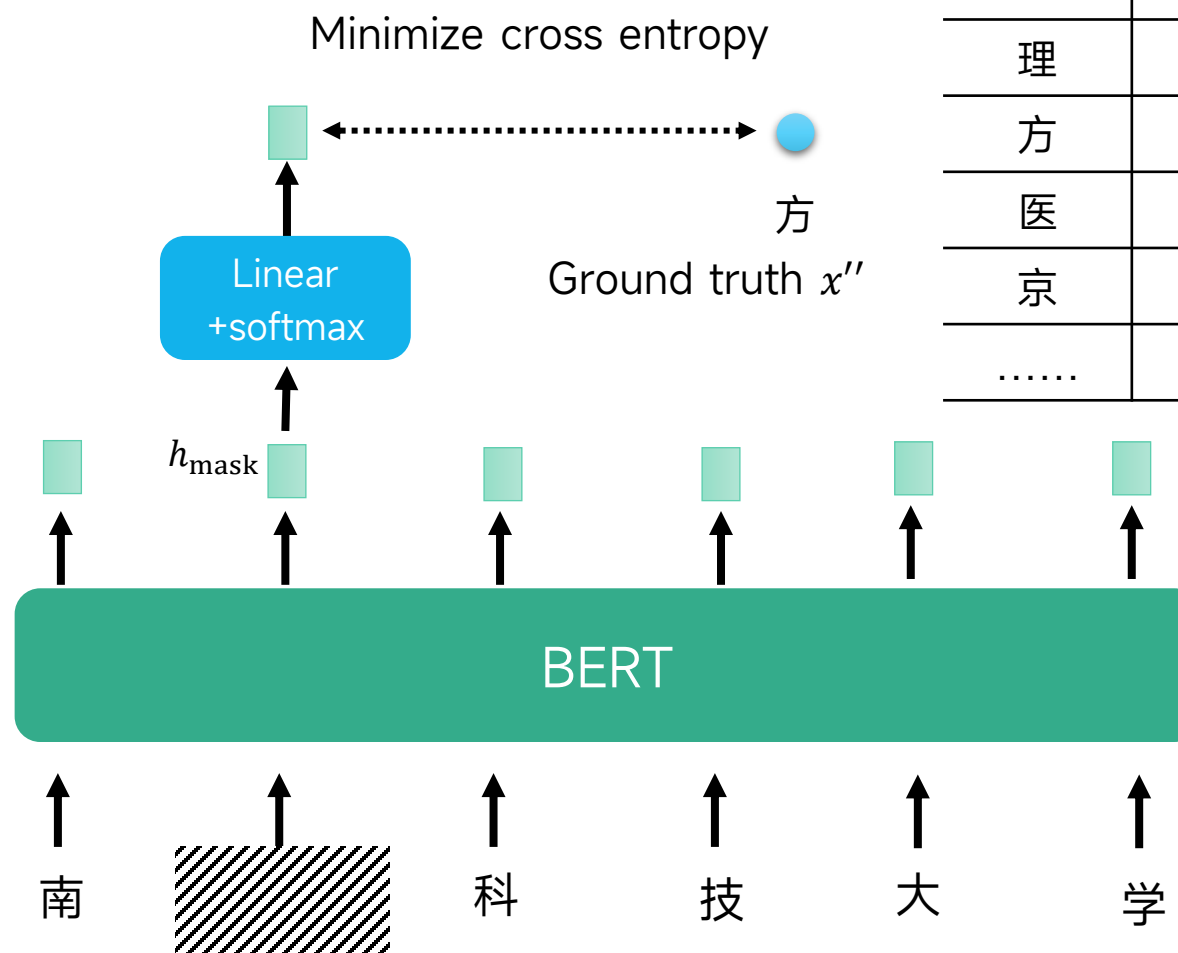


Masked Token Prediction



Transformer Encoder

Input x' (Randomly mask)



Token	Pred	Label
理	0.1	0
方	0.7	1
医	0.1	0
京	0.1	0
.....

MLM: 80-10-10 Corruption



For the 15% predicted words,

- 80% of the time, they replace it with [MASK] token

went to the store → went to the [MASK]

- 10% of the time, they replace it with a random word in the vocabulary

went to the store → went to the running

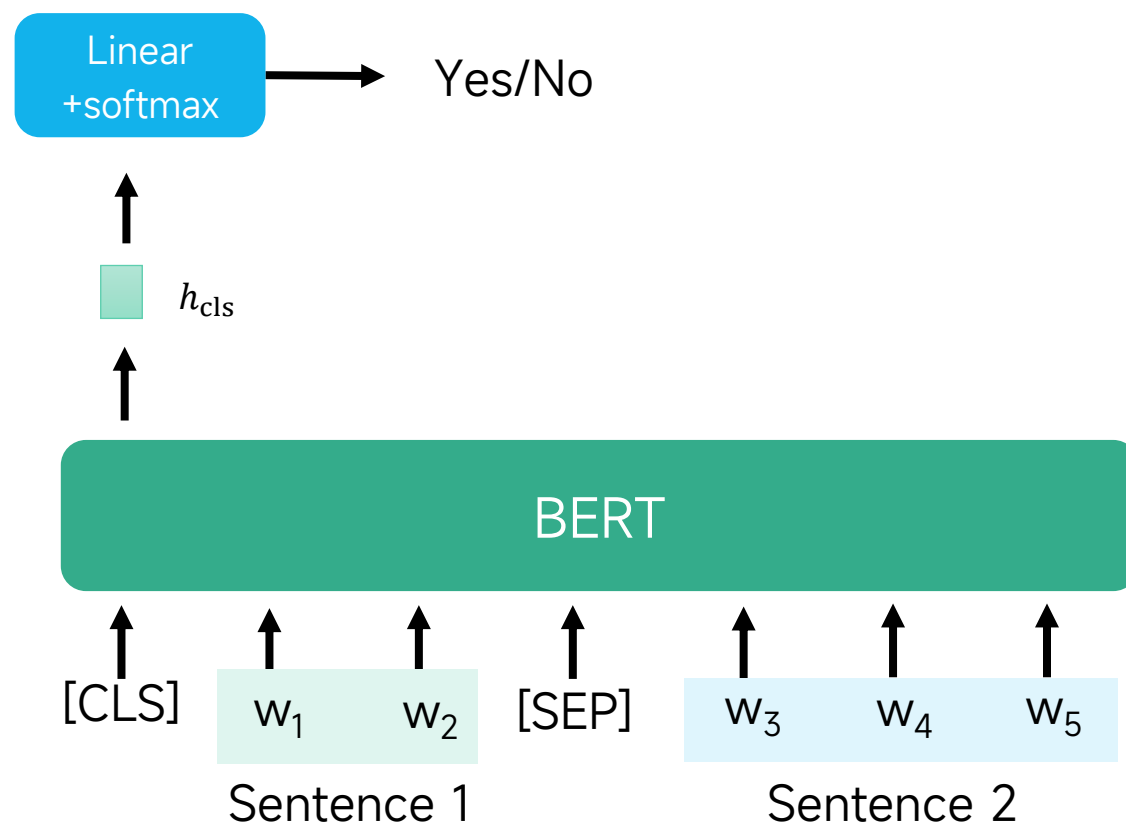
- 10% of the time, they keep it unchanged

went to the store → went to the store

Why?

Because [MASK] tokens are never seen during fine-tuning

Next Sentence Prediction



- This approach is not helpful.
Observed in Roberta
- SOP: Sentence order prediction
Used in ALBERT

<https://arxiv.org/abs/1907.11692>

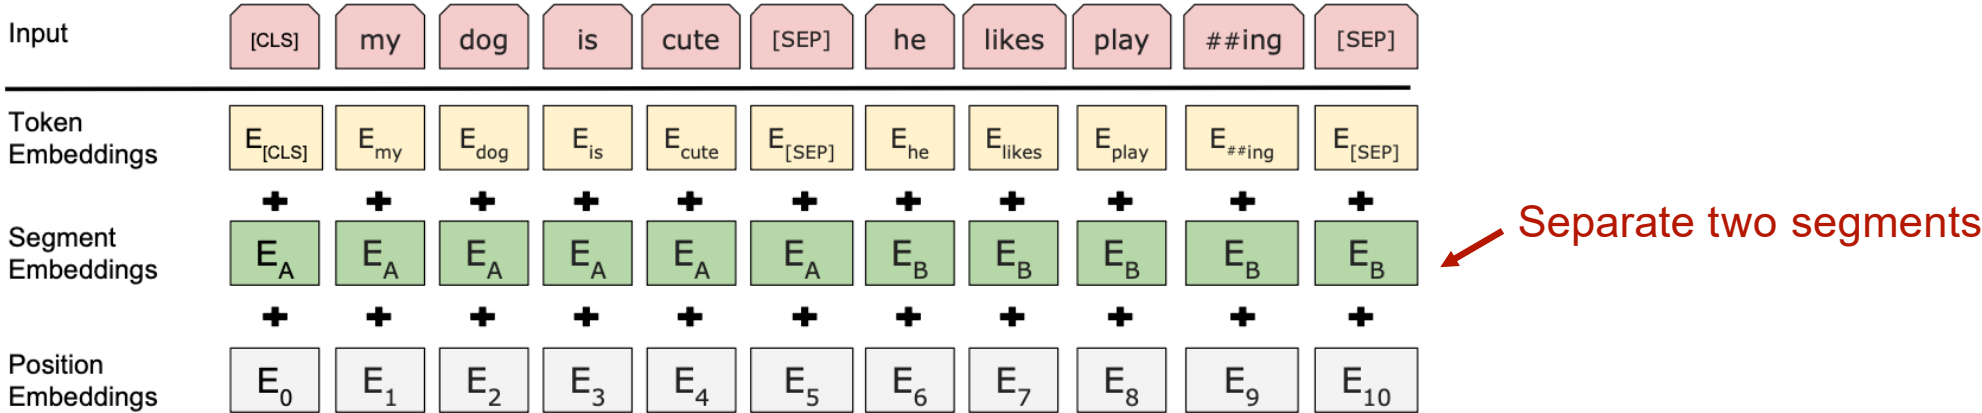
<https://arxiv.org/abs/1909.11942>

BERT Embedding

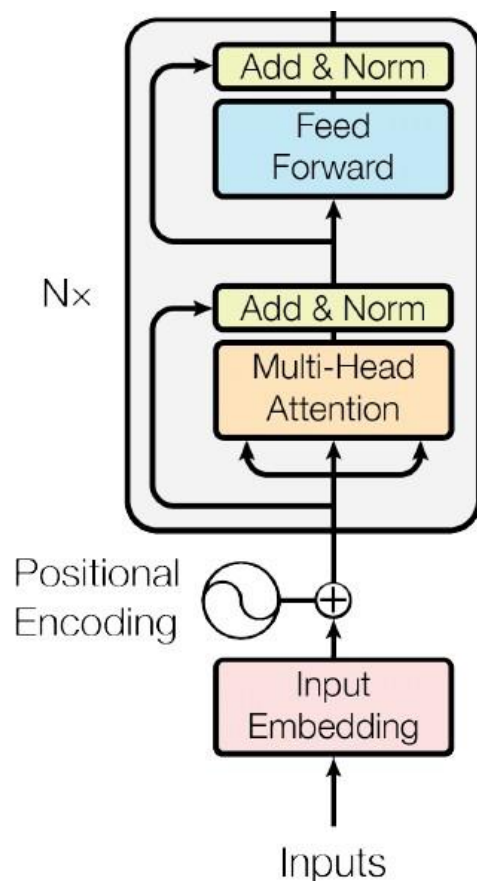
- Vocabulary size: 30,000 wordpieces (common sub-word units) (Wu et al., 2016)



- Input embeddings:



BERT Pretraining



- BERT-base: 12 layers, 768 hidden size, 12 attention heads, 110M parameters
- BERT-large: 24 layers, 1024 hidden size, 16 attention heads, 340M parameters

Same as OpenAI GPT


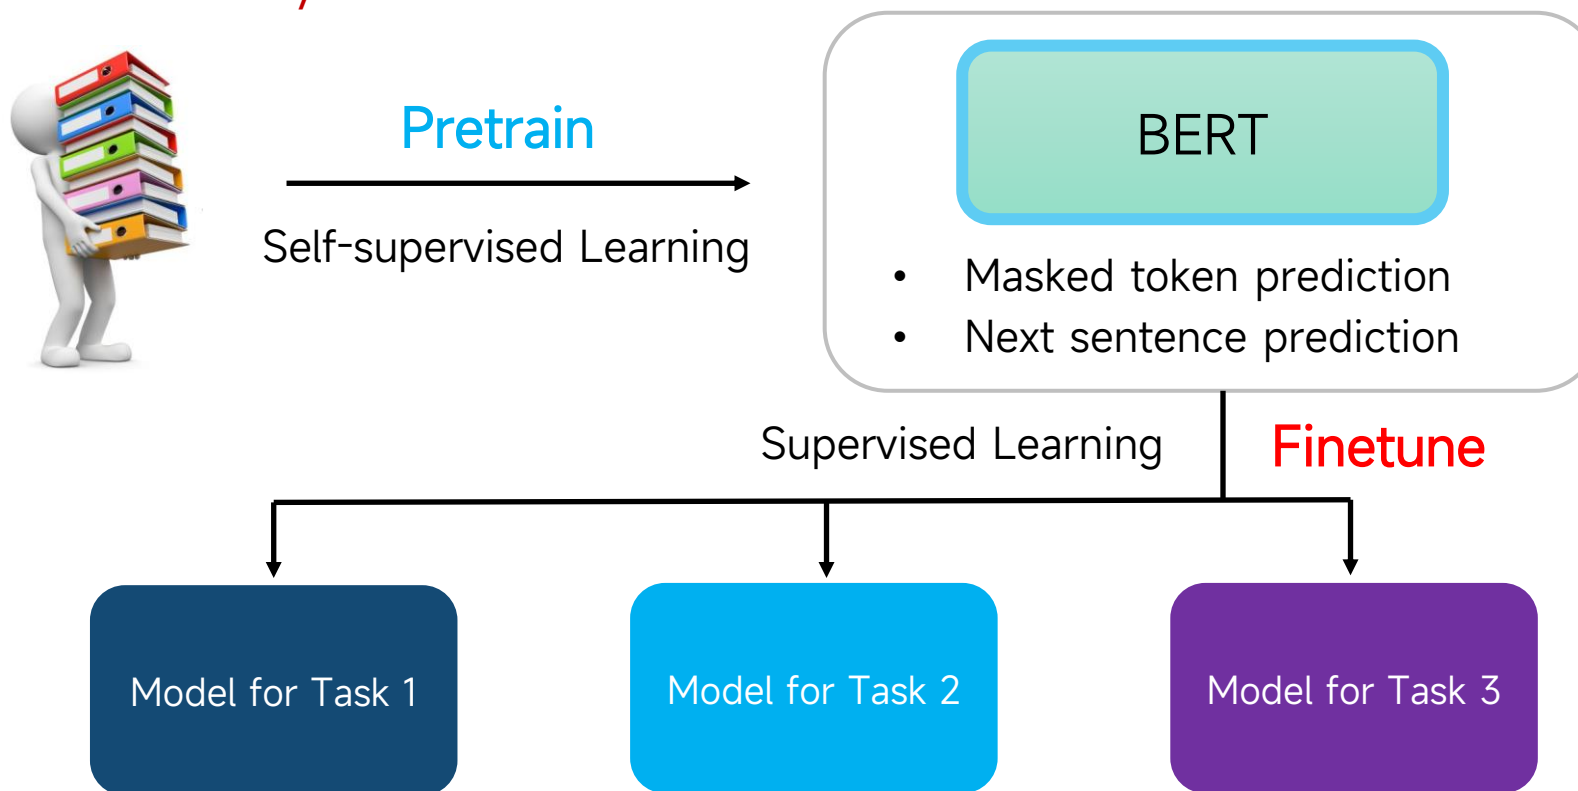
OpenAI GPT was trained on BooksCorpus only!

- Training corpus: Wikipedia (2.5B) + BooksCorpus (0.8B)
- Max sequence size: 512 wordpieces (roughly 256 and 256 for two non-contiguous sequences)
- Trained for 1M steps, batch size 128k

Pretrain-and-Finetune Paradigm



Pretrain once, finetune many times



Embryonic stem cells

Downstream Tasks

- The tasks we care
- We have a little bit of labeled data

Different Downstream Tasks

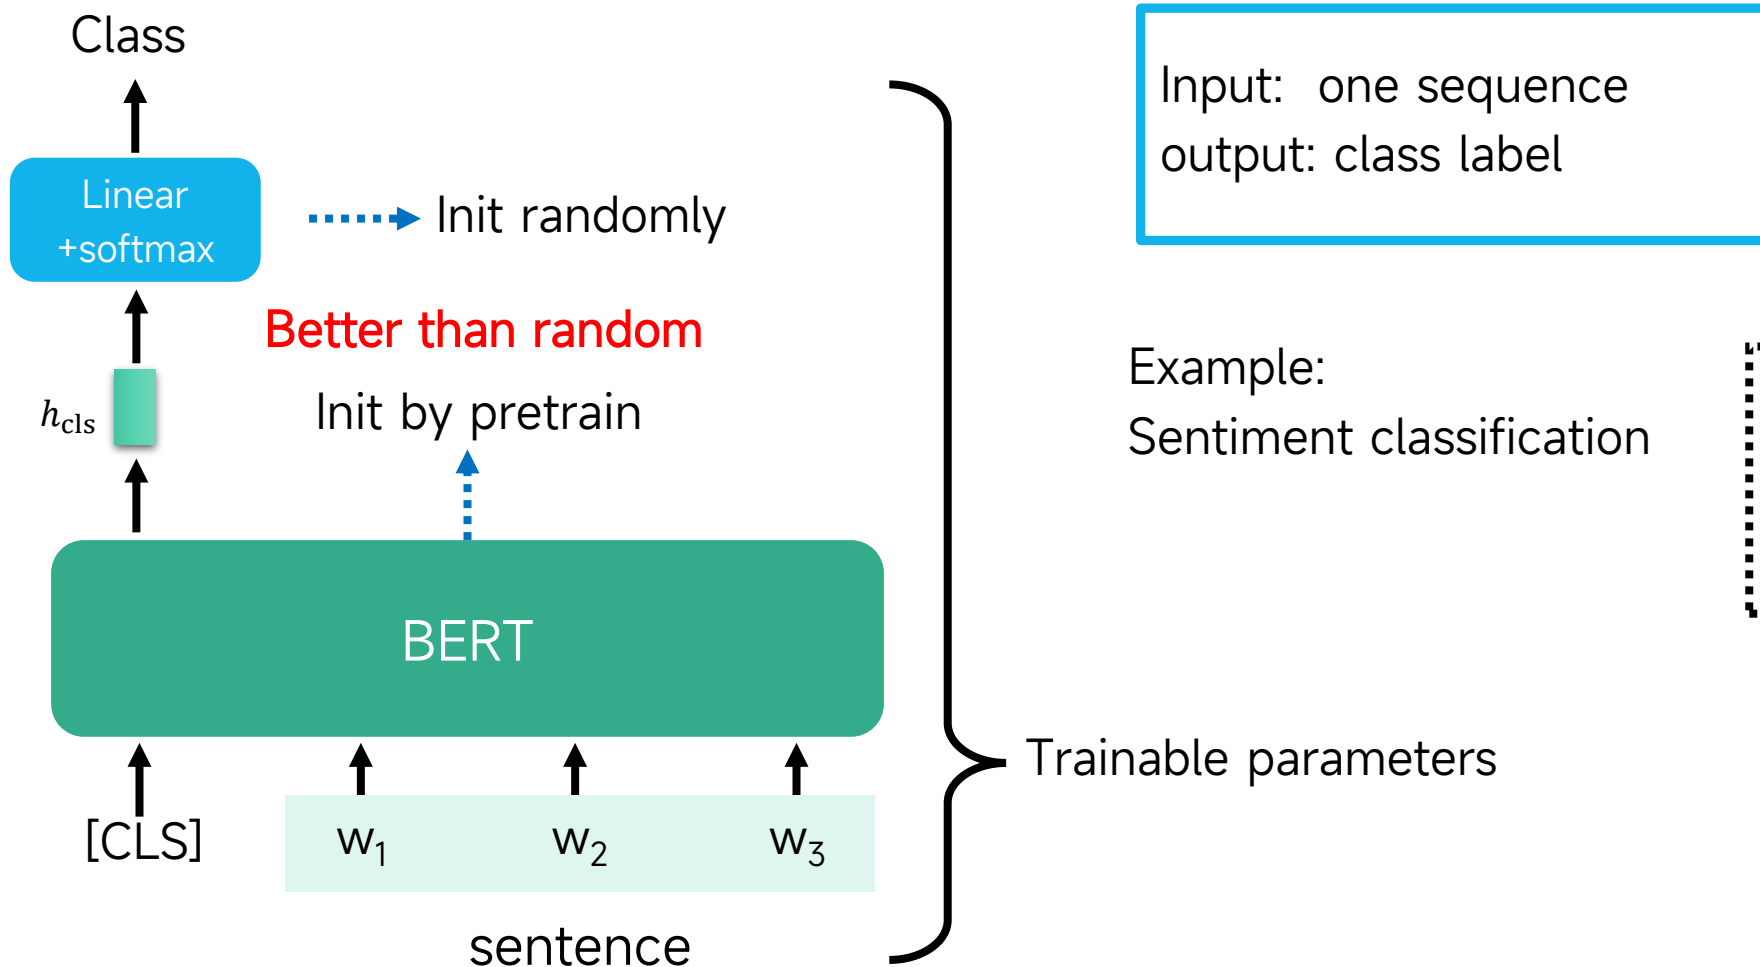


- Text classification
- Sequence labeling
- Matching
- Extractive question answering



Pretrained Encoder

Case 1: Text Classification



The food is good
↓
Positive

Pretrain vs. Random Initialization


(finetune)

(scratch)



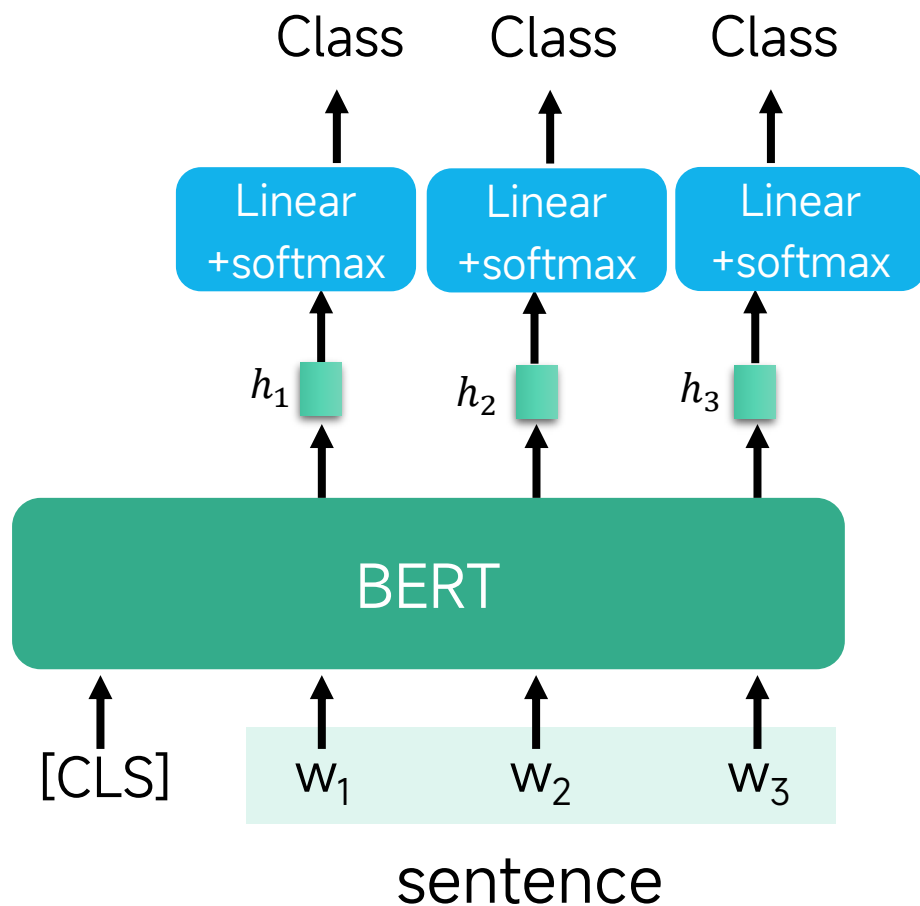
南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Which one is better? Why?



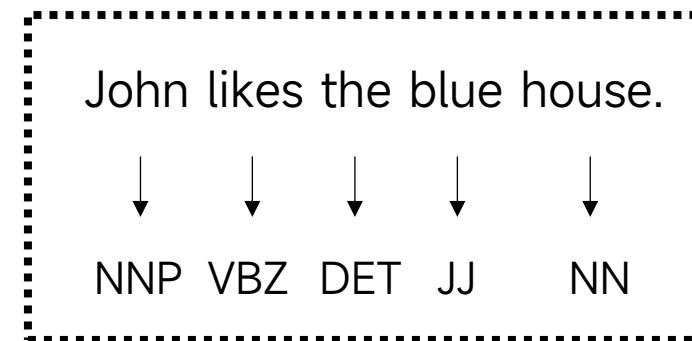
Source of image: <https://arxiv.org/abs/1908.05620>

Case 2: Sequence Labeling

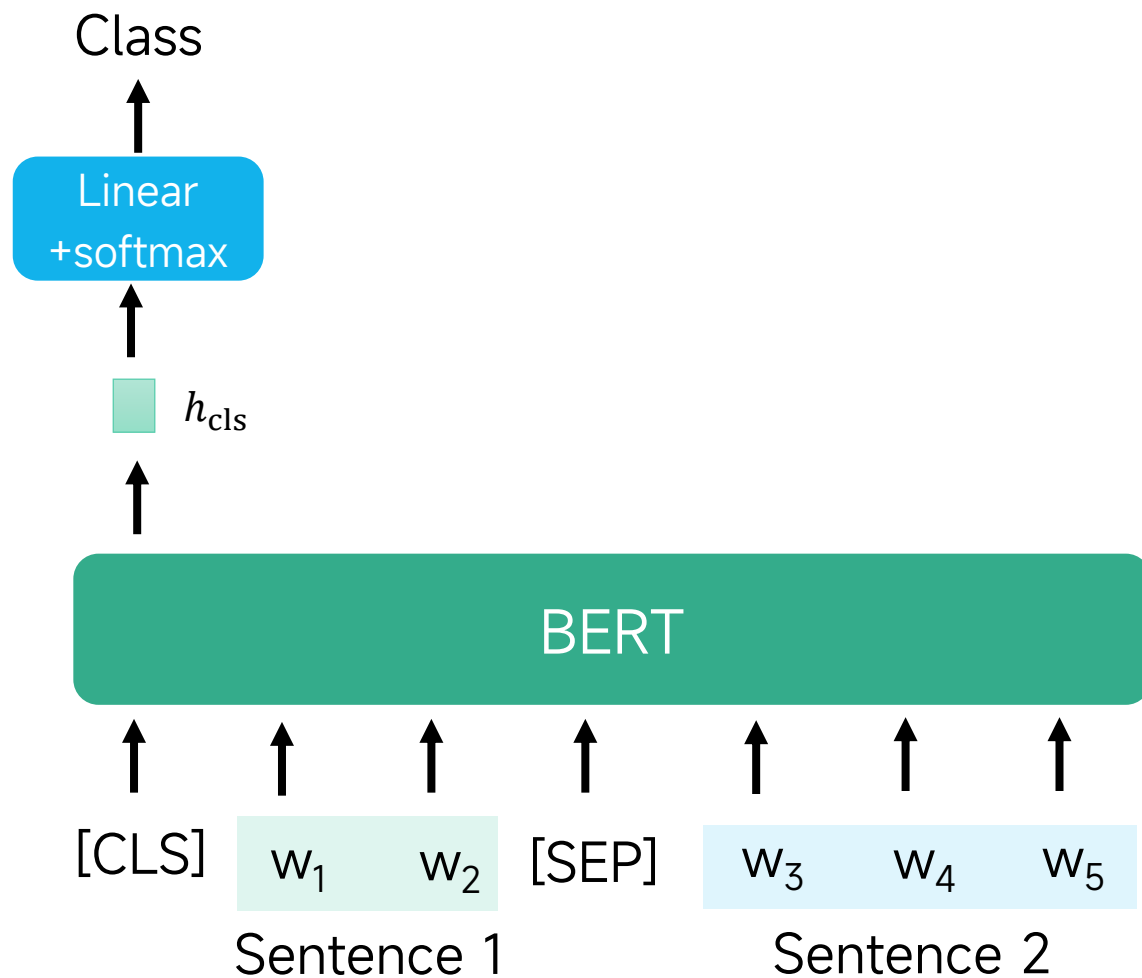


Input: one sequence
Output: one label sequence

Example:
POS tagging



Case 3: Matching



Input: two sequences
output: class label

Example:


Natural Language Inference (NLI)

Premise: A person on a horse jumps over a stone.
Hypothesis: A person is at a diner.
Output: Contradiction

Case 4: Extractive Question Answering



Input: two sequences
Output: answer



Output: two integers (s, e)

Answer. $A = \{d_s, \dots, d_e\}$

In meteorology, precipitation is any product of the condensation of atmospheric water vapor that falls under **gravity**. 17 main forms of precipitation include drizzle, rain, sleet, snow, **graupel** and hail... Precipitation forms as smaller droplets coalesce via collision with other rain drops or ice crystals within a cloud. Short, intense periods of rain in scattered locations are called "showers". 79

What causes precipitation to fall?

gravity $s = 17, e = 17$

What is another main form of precipitation besides drizzle, rain, snow, sleet and hail?

graupel

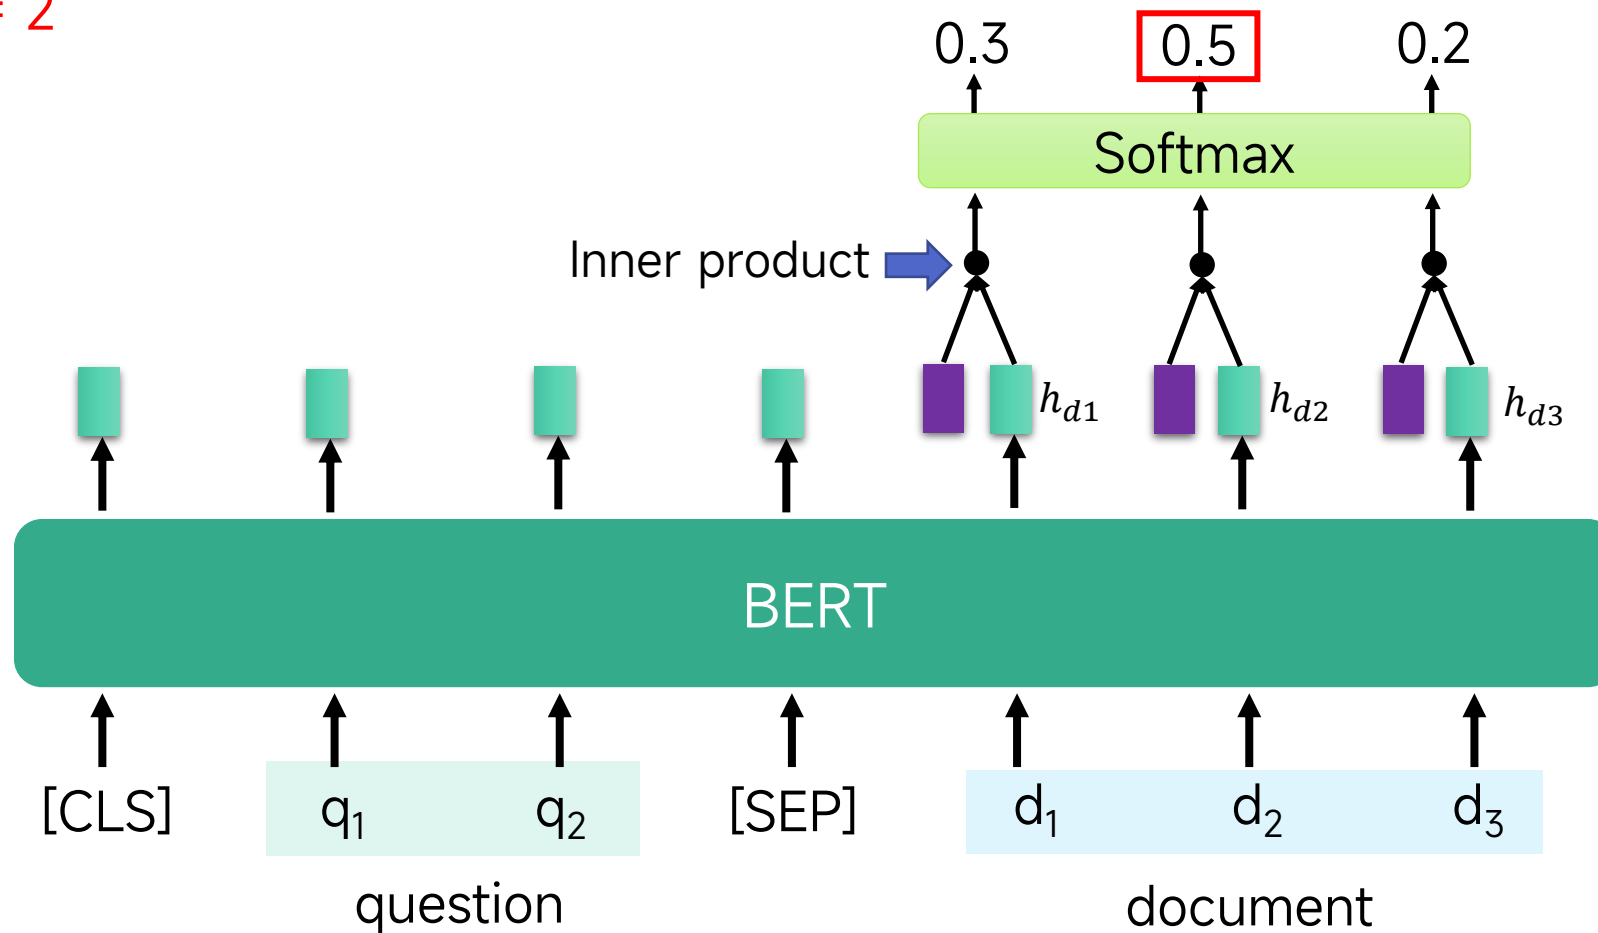
Where do water droplets collide with ice crystals to form precipitation?

within a cloud $s = 77, e = 79$

Case 4: Extractive Question Answering



$s = 2$


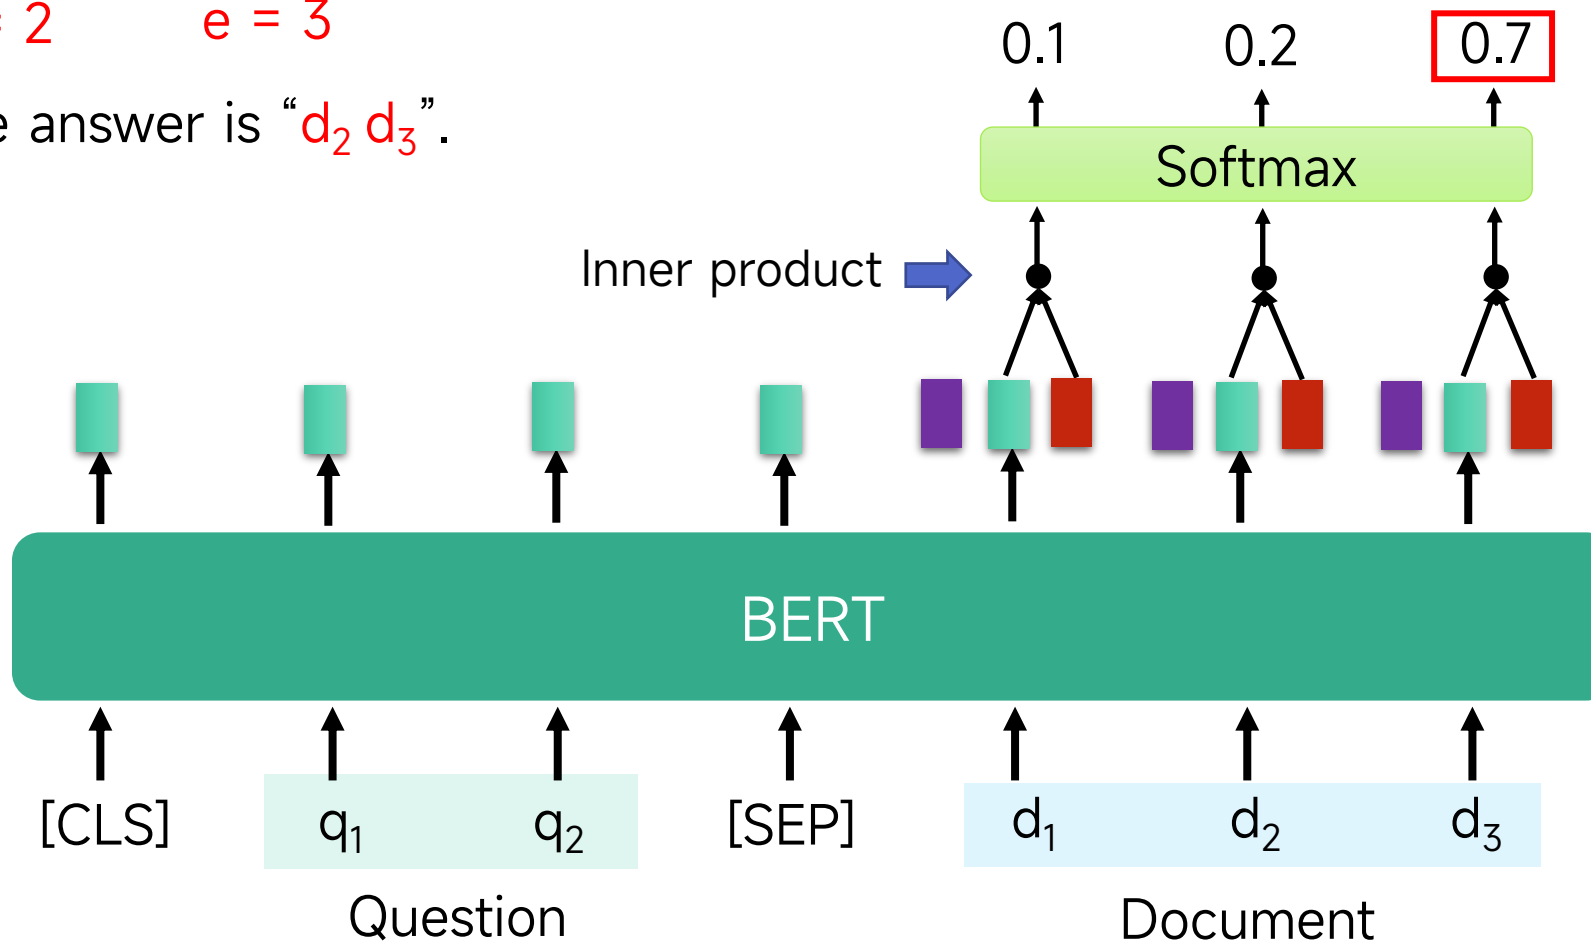


Case 4: Extractive Question Answering



$s = 2$ $e = 3$

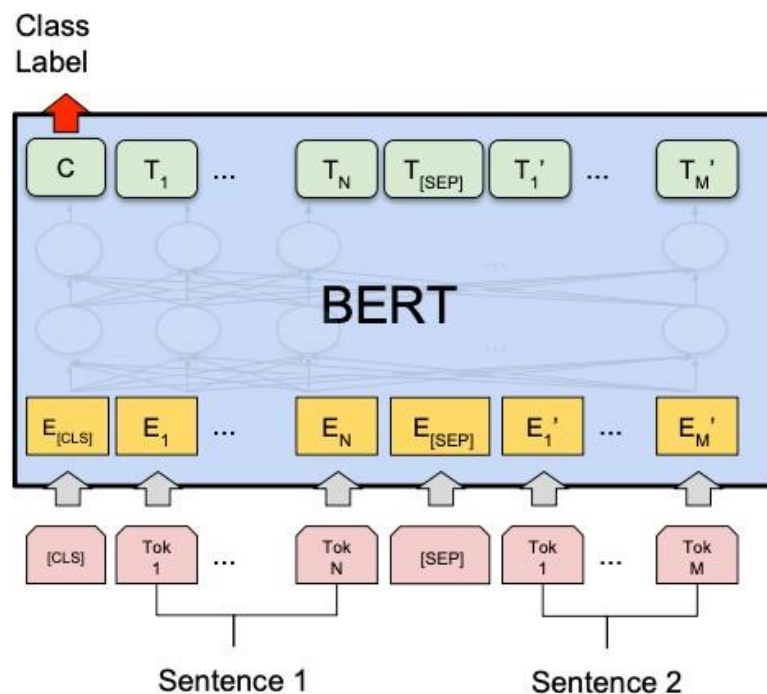
The answer is “ $d_2 d_3$ ”.



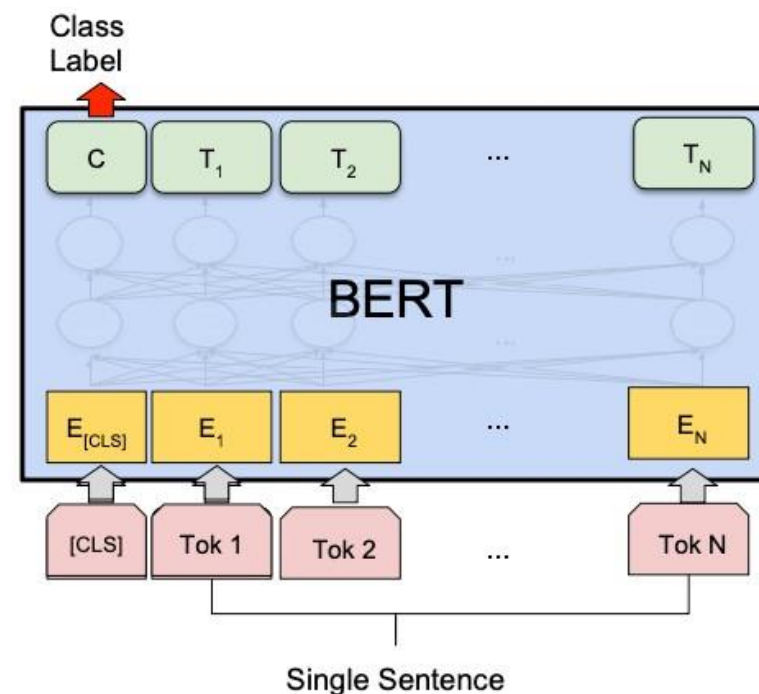
Fine-tuning BERT

“Pretrain once, finetune many times.”

sentence-level tasks



(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG

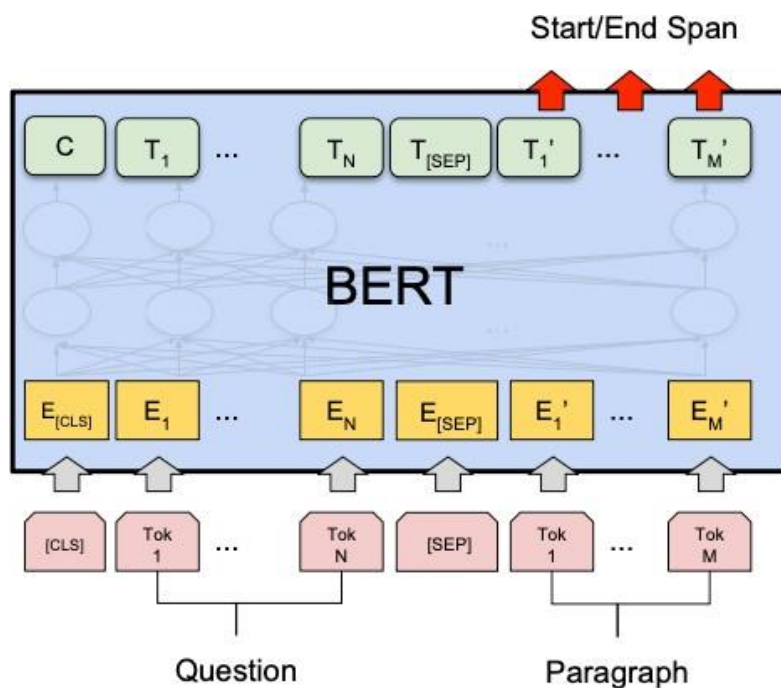


(b) Single Sentence Classification Tasks:
SST-2, CoLA

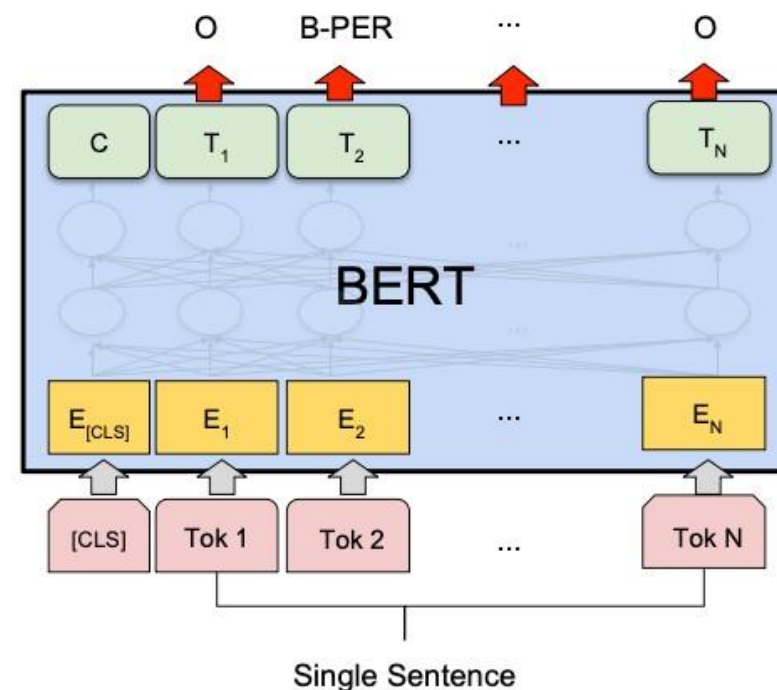
Fine-tuning BERT

“Pretrain once, finetune many times.”

token-level tasks



(c) Question Answering Tasks:
SQuAD v1.1



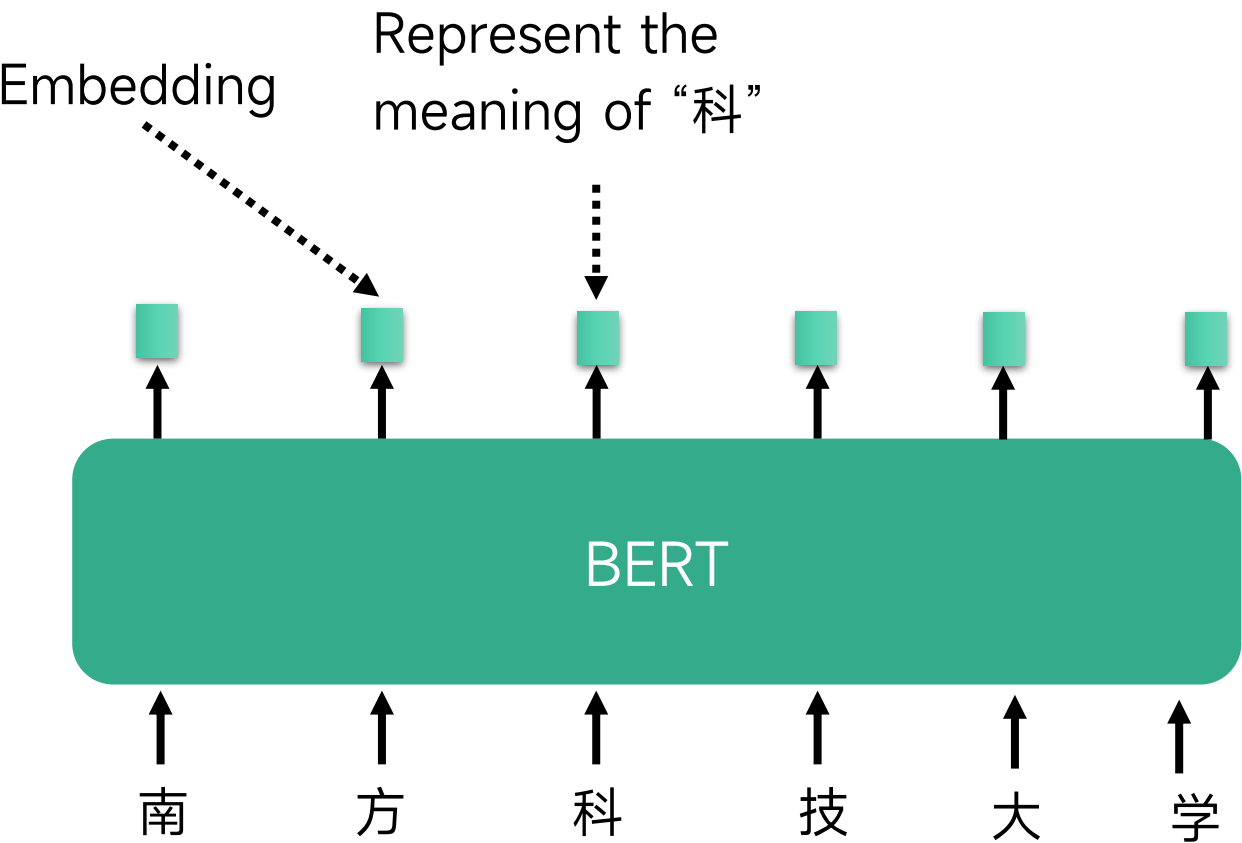
(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

Experimental Results: GLUE




System	MNLI-(m/mm) 392k	QQP 363k	QNLI 108k	SST-2 67k	CoLA 8.5k	STS-B 5.7k	MRPC 3.5k	RTE 2.5k	Average -
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT _{BASE}	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT _{LARGE}	86.7/85.9	72.1	92.7	94.9	60.5	86.5	89.3	70.1	82.1

Why does BERT Work?

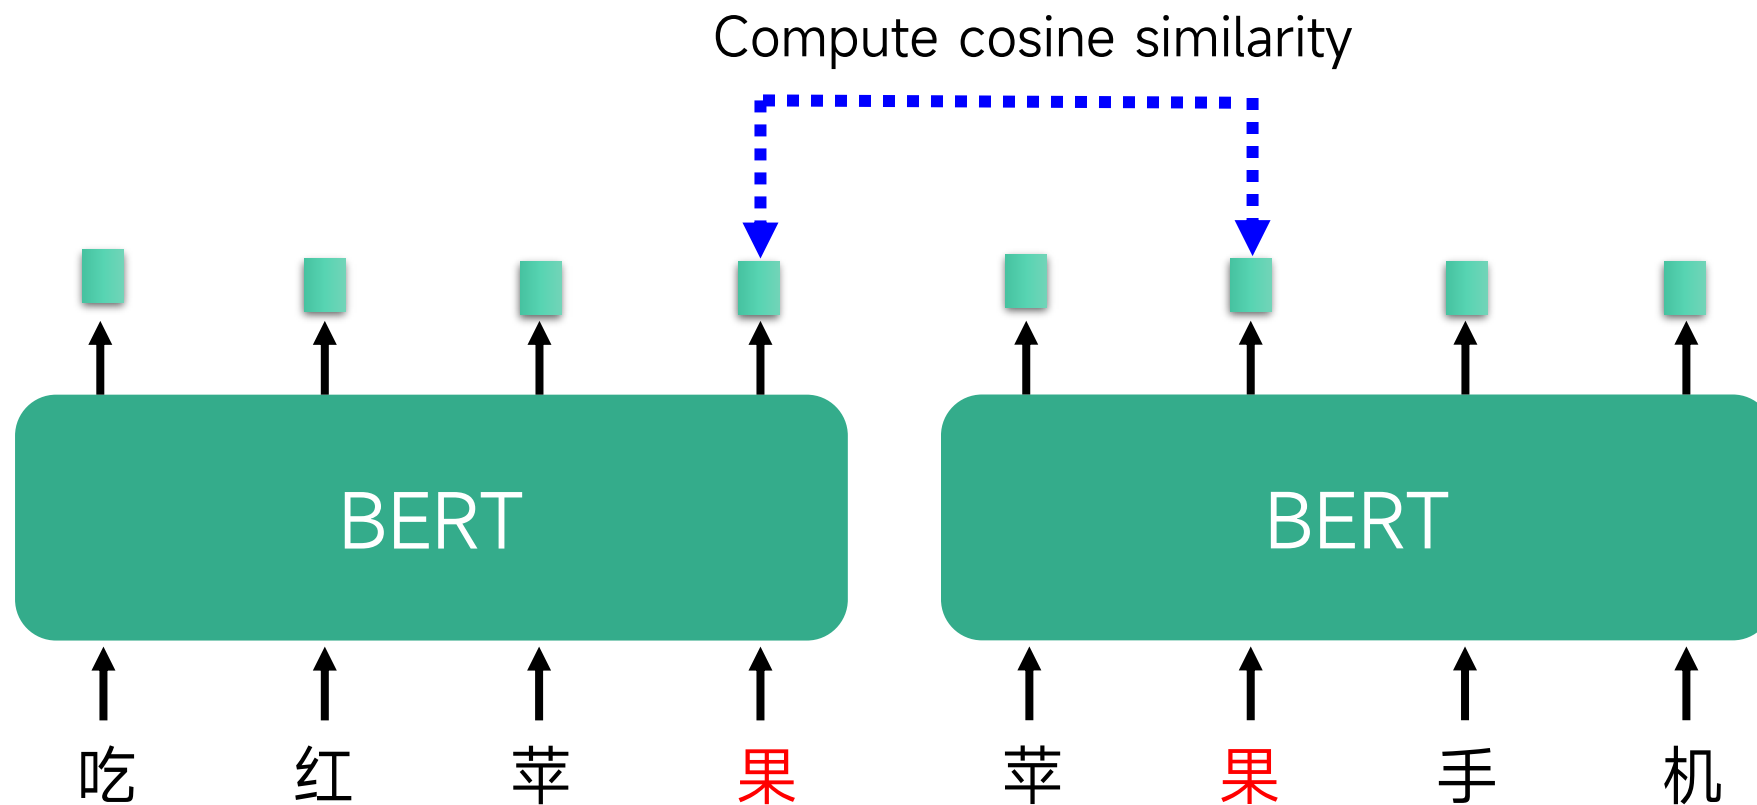


The tokens with similar meaning have similar embedding



Context is considered

Why does BERT Work?



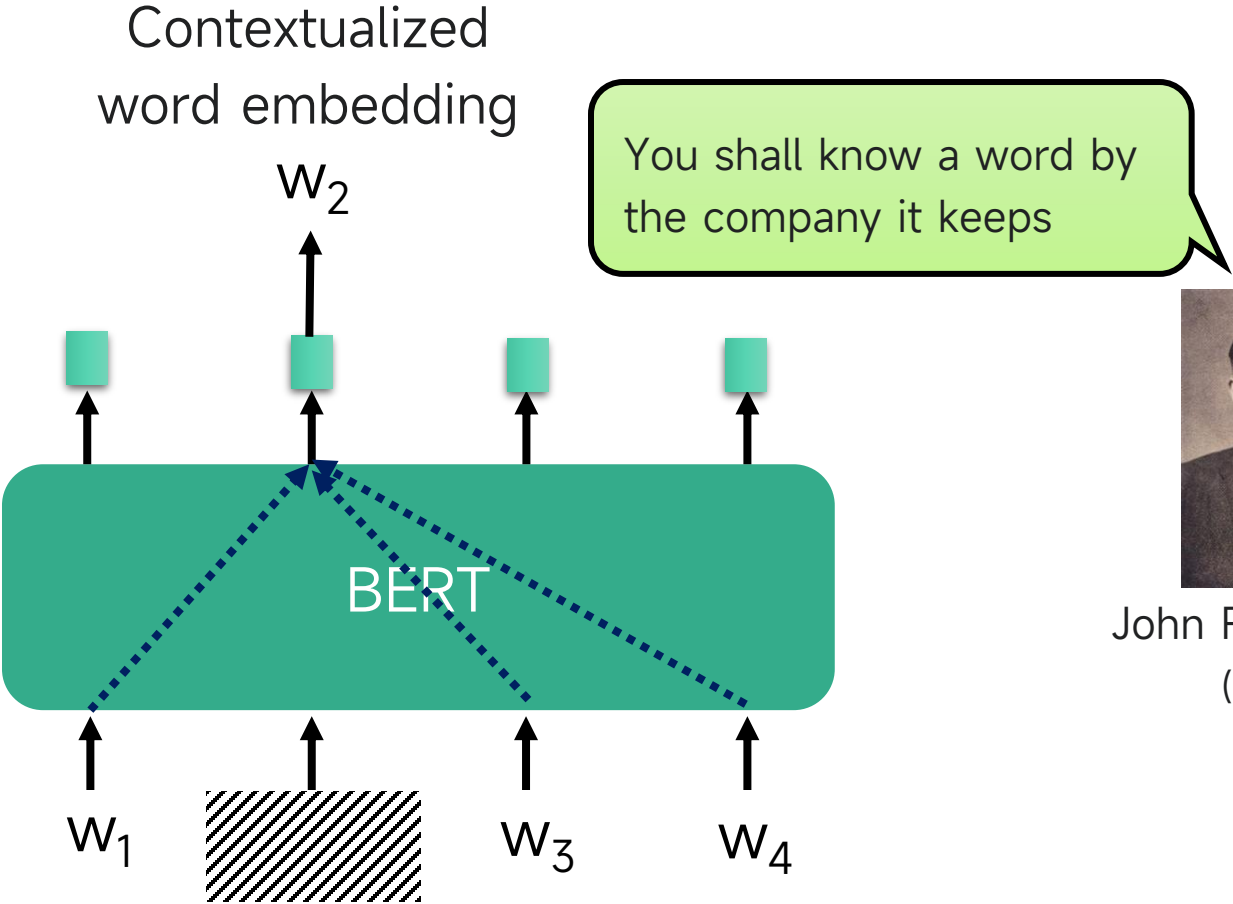
Why does BERT Work?



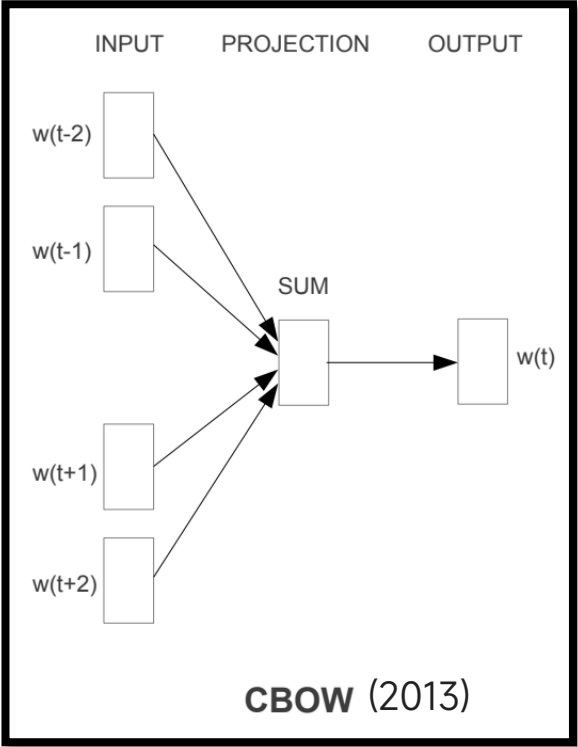
Cosine similarities of BERT embeddings



Why does BERT work?



John Rupert Firth
(1960s)

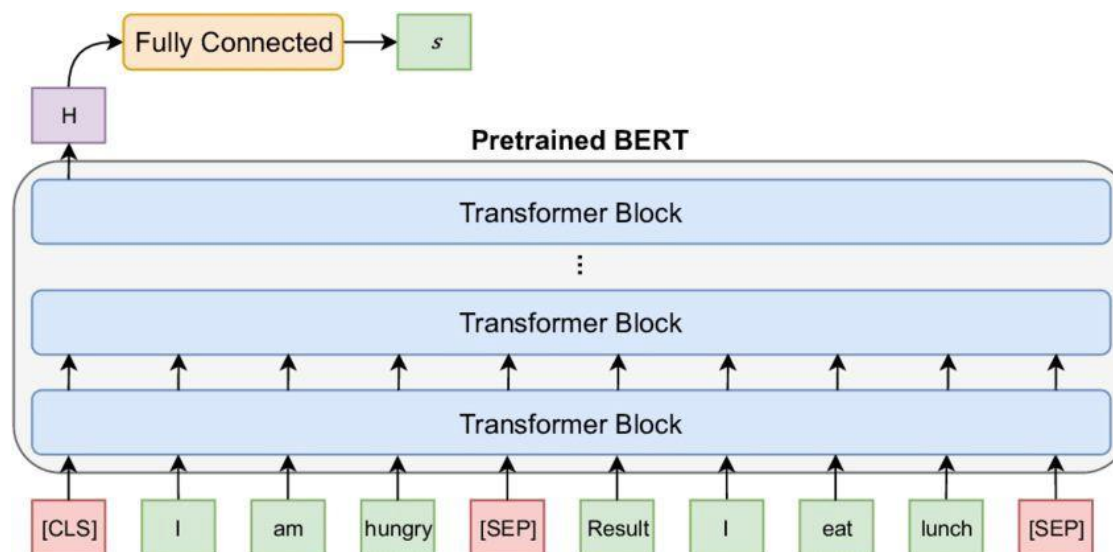


Word2vec

Fully Finetuning



- Pretrain a language model on task
- Attach a small task specific layer
- Fine-tune the weights of full NN by propagating gradients on a downstream task



[Devlin et al. 2019](#)

Parameter-Efficient Finetuning



- With standard fine-tuning, we need to make a new copy of the model for each task
- In the extreme case of a different model per user, we could never store 1000 different full models
- If we fine tuned a subset of the parameters for each task, we could alleviate storage costs
- This is parameter-efficiency

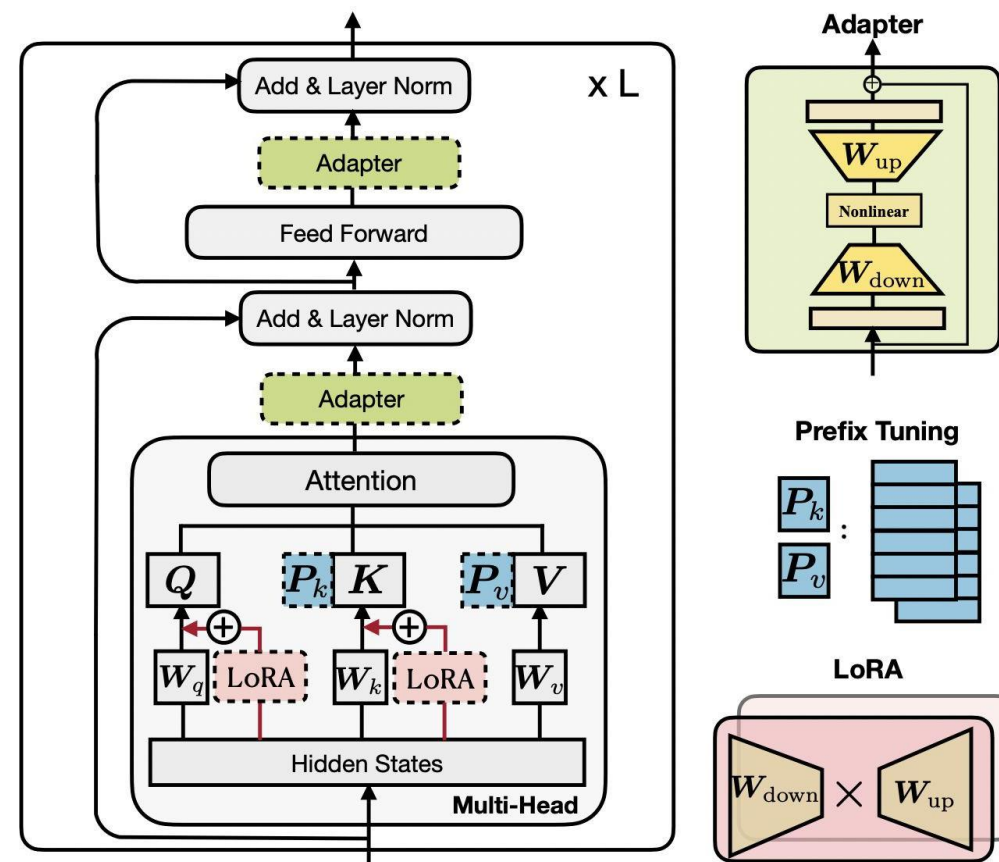
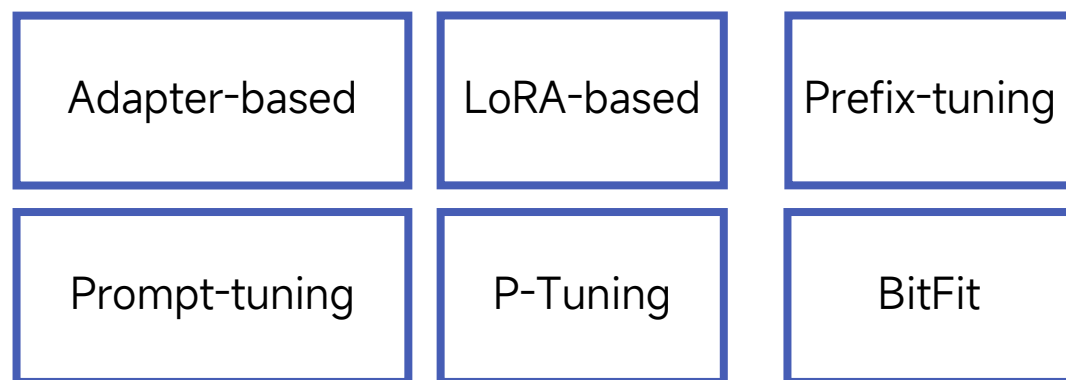


Image: [\(He et al. 2022\)](#)

Tuning the Model

Whether the pretrained models are trained

- Fully finetune (trained)
- Parameter-efficient finetune (fixed)
 - Additional parameters are trained



Pre-train, Prompt, and Predict: A Systematic Survey of Prompting Methods in Natural Language Processing

Pengfei Liu
Carnegie Mellon University
pliu3@cs.cmu.edu


Weizhe Yuan
Carnegie Mellon University
weizhey@cs.cmu.edu

Jinlan Fu
National University of Singapore
jinlanjonna@gmail.com

Zhengbao Jiang
Carnegie Mellon University
zhengbaj@cs.cmu.edu

Hiroaki Hayashi
Carnegie Mellon University
hiroakih@cs.cmu.edu


Graham Neubig
Carnegie Mellon University
gneubig@cs.cmu.edu



LoRA Finetuning



- Low-rank adaptation
- Hypothesizes that the change of weights during model tuning has a low intrinsic rank
- Optimize the low-rank decomposition for the change of original weight matrices in the self-attention modules.



$$\text{LoRA } (W = W_0 + \gamma BA)$$
$$B \in \mathbb{R}^{d \times r}, A \in \mathbb{R}^{r \times d}$$

[Parameter-Efficient LLM Finetuning With Low-Rank Adaptation \(LoRA\) - Lightning AI](#)

LoRA Finetuning

- Code illustration

```
input_dim = 768 # e.g., the hidden size of the pre-trained model
output_dim = 768 # e.g., the output size of the layer
rank = 8 # The rank 'r' for the low-rank adaptation


W = ... # from pretrained network with shape input_dim x output_dim

W_A = nn.Parameter(torch.empty(input_dim, rank)) # LoRA weight A
W_B = nn.Parameter(torch.empty(rank, output_dim)) # LoRA weight B

# Initialization of LoRA weights
nn.init.kaiming_uniform_(W_A, a=math.sqrt(5))
nn.init.zeros_(W_B)

def regular_forward_matmul(x, W):
    h = x @ W
    return h

def lora_forward_matmul(x, W, W_A, W_B):
    h = x @ W # regular matrix multiplication
    h += x @ (W_A @ W_B)*alpha # use scaled LoRA weights
    return h
```



LoRA Finetuning



- Code example

```
import peft
from transformers import AutoModelForCausalLM, AutoTokenizer
from peft import LoraConfig, get_peft_model, PeftModel

lora_config = LoraConfig(
    r=4, #As bigger the R bigger the parameters to train.
    lora_alpha=1, # a scaling factor that adjusts the magnitude of the weight
    matrix. Usually set to 1
    target_modules=["query_key_value"], #You can obtain a list of target modules
    in the URL above.
    lora_dropout=0.05, #Helps to avoid Overfitting.
    bias="lora_only", # this specifies if the bias parameter should be trained.
    task_type="CAUSAL_LM"
)

model_name = "Qwen/Qwen2.5-1.5B-Instruct"
tokenizer = AutoTokenizer.from_pretrained(model_name)
foundation_model = AutoModelForCausalLM.from_pretrained(model_name)

peft_model = get_peft_model(foundation_model, lora_config)
print(peft_model.print_trainable_parameters())
```

LoRA Finetuning

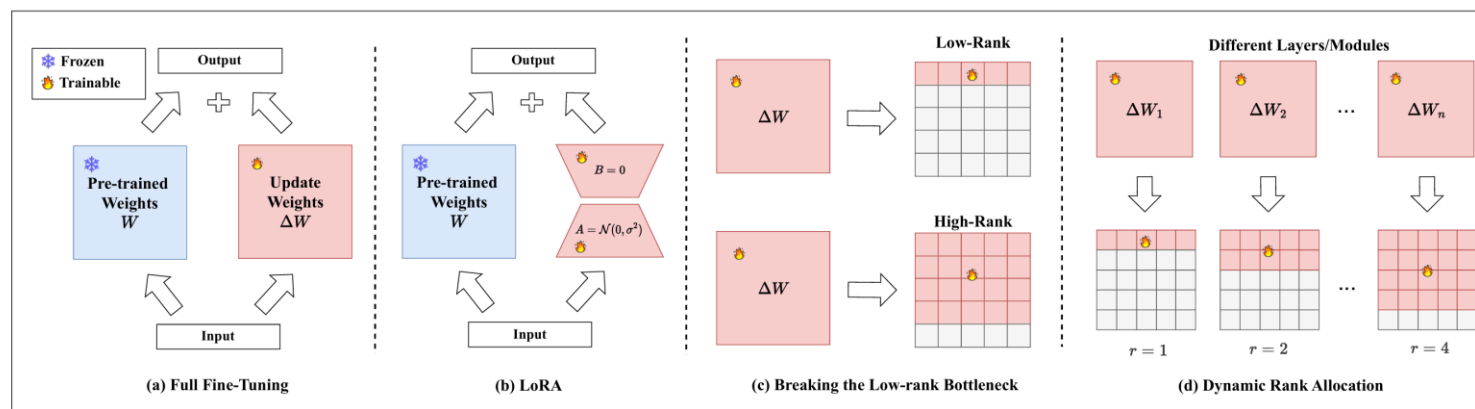


Fig. 2 An illustration of full fine-tuning (a), LoRA (b) and its variants for improving downstream adaptation, which includes breaking the low-rank bottleneck (c) and dynamic rank allocation (d).

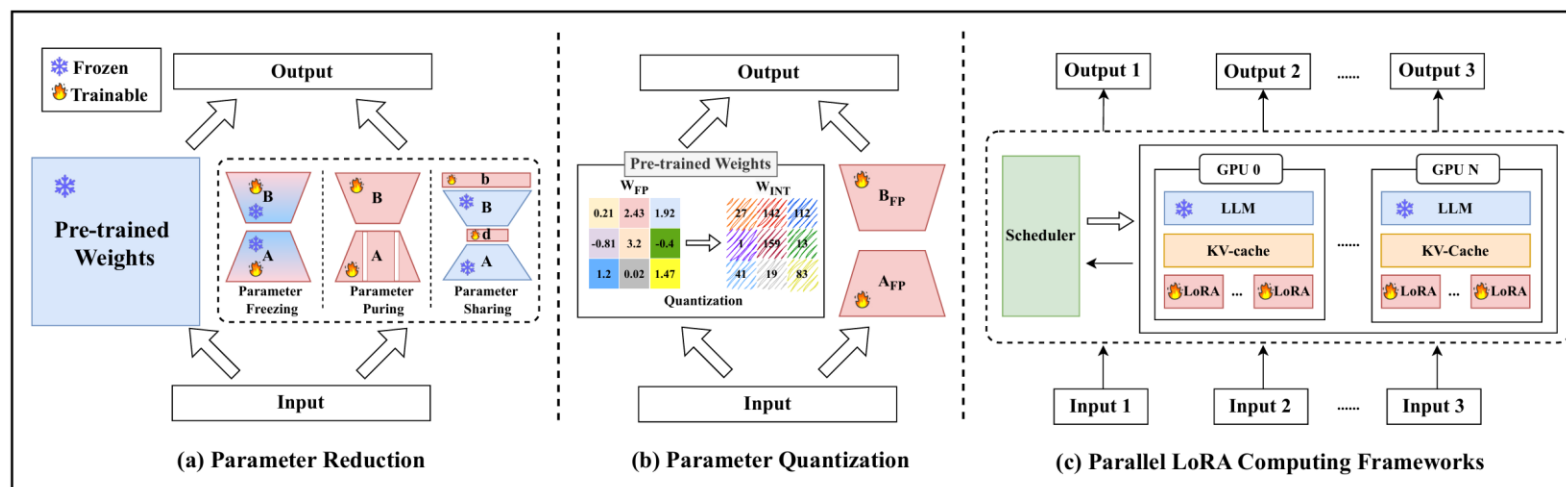


Fig. 4 An illustration of efficiency improving methods.

What Happened after BERT?

Lots of people are trying to understand what BERT has learned and how it works

A Primer in BERTology: What We Know About How BERT Works

Anna Rogers

Center for Social Data Science
University of Copenhagen
arogers@sodas.ku.dk

Olga Kovaleva

Dept. of Computer Science
University of Massachusetts Lowell
okovalev@cs.uml.edu

Anna Rumshisky

Dept. of Computer Science
University of Massachusetts Lowell
arum@cs.uml.edu

- Syntactic knowledge, semantic knowledge, world knowledge...
- How to mask, what to mask, where to mask, alternatives to masking..

What Happened after BERT?

- Models that handle long contexts (much more than 512 tokens)
 - Longformer, Big Bird, ...
- Multilingual BERT
 - mBERT: Trained single model on 104 languages from Wikipedia. Shared 110k WordPiece vocabulary
 - XLM-R: 100 languages with 250k sentencepiece vocabulary.
- BERT extended to different domains
 - SciBERT, BioBERT, FinBERT, ClinicalBERT, ...
- Making BERT smaller to use
 - DistillBERT, TinyBERT, ...

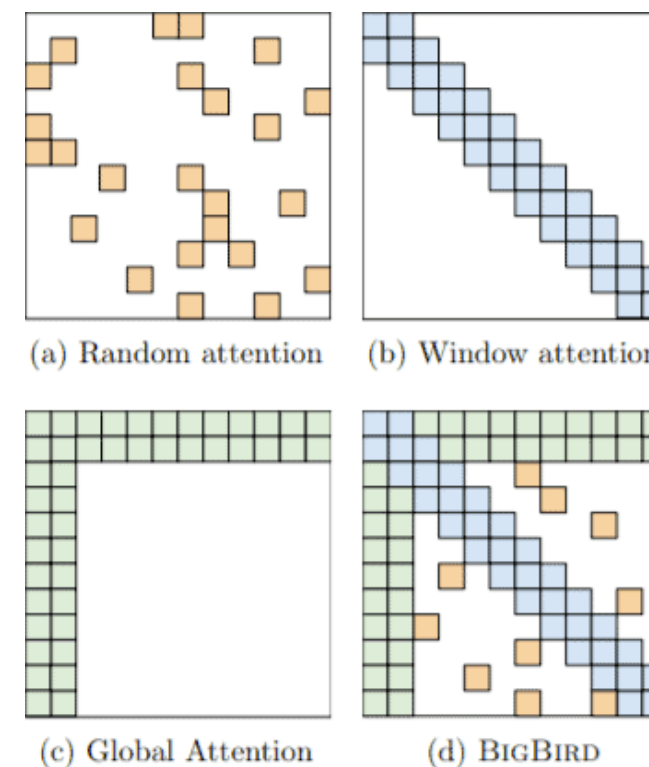


Image from the original paper

What Happened after BERT?

- RoBERTa (Liu et al., 2019)
 - Trained on 10x data & longer, no NSP
 - Much stronger performance than BERT (e.g., 94.6 vs 90.9 on SQuAD)
 - Still one of the most popular models to date
- ALBERT (Lan et al., 2020)
 - Increasing model sizes by sharing model parameters across layers
 - Less storage, much stronger performance but runs slower..
- ModernBERT (Warner et al., 2024)
 - Train on 2 trillion tokens with a native 8192 sequence length
 - RoPE, Pre-Norm, GeGLU

[AnswerDotAI/ModernBERT: Bringing BERT into modernity via both architecture changes and scaling](#)

T5 Model



- T5 refers to “Text-to-Text Transfer Transformer”

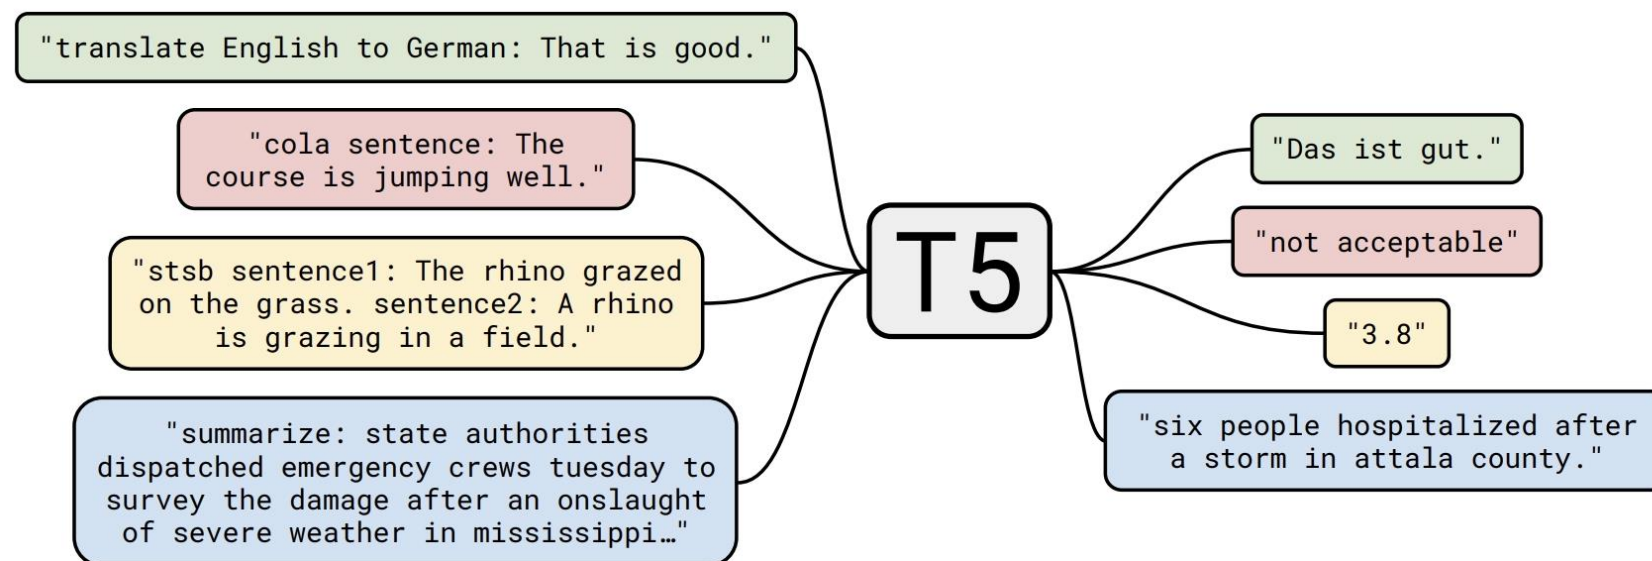


Figure 1: A diagram of our text-to-text framework. Every task we consider—including

C4 Dataset

- Colossal Clean Crawled Corpus (750 GB)
- Clean the data from Common Crawl (20TB data each month)

- We only retained lines that ended in a terminal punctuation mark (i.e. a period, exclamation mark, question mark, or end quotation mark).
- We discarded any page with fewer than 3 sentences and only retained lines that contained at least 5 words.
- We removed any page that contained any word on the “List of Dirty, Naughty, Obscene or Otherwise Bad Words”.⁶
- Many of the scraped pages contained warnings stating that Javascript should be enabled so we removed any line with the word Javascript.
- Some pages had placeholder “lorem ipsum” text; we removed any page where the phrase “lorem ipsum” appeared.
- Some pages inadvertently contained code. Since the curly bracket “{” appears in many programming languages (such as Javascript, widely used on the web) but not in natural text, we removed any pages that contained a curly bracket.
- Since some of the scraped pages were sourced from Wikipedia and had citation markers (e.g. [1], [citation needed], etc.), we removed any such markers.
- Many pages had boilerplate policy notices, so we removed any lines containing the strings “terms of use”, “privacy policy”, “cookie policy”, “uses cookies”, “use of cookies”, or “use cookies”.
- To deduplicate the data set, we discarded all but one of any three-sentence span occurring more than once in the data set.

T5 Input and Output Format



- Train a single model on the diverse set of tasks
- Cast all of the tasks we consider into a “text-to-text” format
 - “[Task-specific prefix]: [Input text]” -> “[output text]”
 - Add a task-specific (text) prefix to the original input sequence
 - For translation, add the sequence “translate English to German: That is good.”
 - For classification, simply predicts a single word corresponding to the target label
- The choice of text prefix used for a given task is essentially a hyperparameter
- Provides a consistent training objective both for pre-training and fine-tuning

T5 Pretraining



- Pretrain a standard Transformer using a simple denoising objective and then separately fine-tune on each of our downstream tasks
- A model with about 220 million parameters. Roughly twice the number of parameters of Bert-base
- Use an “inverse square root” learning rate schedule:

$$\frac{L_r}{\sqrt{\max(n, k)}}$$

- Where n is the current training iteration and k is the number of warm-up steps (set to 10^4)
- Sets a constant learning rate of 0.01 for the first 10^4 steps, then exponentially decays the learning rate until pre-training is over

T5 Pretraining

Objective	Inputs	Targets
Prefix language modeling	Thank you for inviting	me to your party last week .
BERT-style	Thank you <M> <M> me to your party apple week .	(original text)
Deshuffling	party me for your to . last fun you inviting week Thank	(original text)
I.i.d. noise, mask tokens	Thank you <M> <M> me to your party <M> week .	(original text)
I.i.d. noise, replace spans	Thank you <X> me to your party <Y> week .	<X> for inviting <Y> last <Z>
I.i.d. noise, drop tokens	Thank you me to your party week .	for inviting last
Random spans	Thank you <X> to <Y> week .	<X> for inviting me <Y> your party last <Z>

Original text

Thank you for inviting me to your party last week.


Inputs

Thank you <X> me to your party <Y> week.

Targets


<X> for inviting <Y> last <Z>

- Different Attention Mask Patterns




3: Matrices representing different attention mask patterns. The input and output of the self-attention mechanism are denoted x and y respectively. A dark cell at row i and column j indicates that the self-attention mechanism is allowed to attend to input element j at output timestep i . A light cell indicates that the


T5 Pretraining



Language model



Prefix LM



Performance of Different Variants

Architecture	Objective	Params	Cost	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ Encoder-decoder	Denoising	$2P$	M	83.28	19.24	80.88	71.36	26.98	39.82	27.65
Enc-dec, shared	Denoising	P	M	82.81	18.78	80.63	70.73	26.72	39.03	27.46
Enc-dec, 6 layers	Denoising	P	$M/2$	80.88	18.97	77.59	68.42	26.38	38.40	26.95
Language model	Denoising	P	M	74.70	17.93	61.14	55.02	25.09	35.28	25.86
Prefix LM	Denoising	P	M	81.82	18.61	78.94	68.11	26.43	37.98	27.39
Encoder-decoder	LM	$2P$	M	79.56	18.59	76.02	64.29	26.27	39.17	26.86
Enc-dec, shared	LM	P	M	79.60	18.13	76.35	63.50	26.62	39.17	27.05
Enc-dec, 6 layers	LM	P	$M/2$	78.67	18.26	75.32	64.06	26.13	38.42	26.89
Language model	LM	P	M	73.78	17.54	53.81	56.51	25.23	34.31	25.38
Prefix LM	LM	P	M	79.68	17.84	76.87	64.86	26.28	37.51	26.76

1. Sharing parameters in encoder and decoder models perform nearly as well as the baseline.
2. Halving the number of layers in encoder and decoder hurts the performance.
3. Performance of Encoder and Decoder with shared parameters is better than decoder only LM and prefix LM.

Further Reading



- [The Transformer Family Version 2.0](#)
- [乘风破浪的PTM：两年来预训练模型的技术进展](#)
- [通向AGI之路：大型语言模型（LLM）技术精要](#)



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Thank you